

# Course Overview

15-213: Introduction to Computer Systems  
1<sup>st</sup> Lecture, Sep. 1, 2015

## Instructors:

Randal E. Bryant and David R. O'Hallaron

*The course that gives CMU its “Zip”!*

# Overview

- **Course theme**
- **Five realities**
- **How the course fits into the CS/ECE curriculum**
- **Academic integrity**

# Course Theme:

## Abstraction Is Good But Don't Forget Reality

### ■ Most CS and CE courses emphasize abstraction

- Abstract data types
- Asymptotic analysis

### ■ These abstractions have limits

- Especially in the presence of bugs
- Need to understand details of underlying implementations

### ■ Useful outcomes from taking 213


- Become more effective programmers
  - Able to find and eliminate bugs efficiently
  - Able to understand and tune for program performance
- Prepare for later “systems” classes in CS & ECE
  - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems, Storage Systems, etc.

# Great Reality #1:

## Ints are not Integers, Floats are not Reals

### ■ Example 1: Is $x^2 \geq 0$ ?

■ Float's: Yes!



■ Int's:

- $40000 * 40000 \approx 1600000000$
- $50000 * 50000 \approx ?$

### ■ Example 2: Is $(x + y) + z = x + (y + z)$ ?

■ Unsigned & Signed Int's: Yes!

■ Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

# Computer Arithmetic

## ■ Does not generate random values

- Arithmetic operations have important mathematical properties

## ■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
  - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
  - Monotonicity, values of signs

## ■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

# Great Reality #2:

## You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
  - Compilers are much better & more patient than you are
- **But: Understanding assembly is key to machine-level execution model**
  - Behavior of programs in presence of bugs
    - High-level language models break down
  - Tuning program performance
    - Understand optimizations done / not done by the compiler
    - Understanding sources of program inefficiency
  - Implementing system software
    - Compiler has machine code as target
    - Operating systems must manage process state
  - Creating / fighting malware
    - x86 assembly is the language of choice!

# Great Reality #3: Memory Matters

## Random Access Memory Is an Unphysical Abstraction

### ■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated

### ■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

### ■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

# Memory Referencing Bug Example

```
typedef struct {  
    int a[2];  
    double d;  
} struct_t;  
  
double fun(int i) {  
    volatile struct_t s;  
    s.d = 3.14;  
    s.a[i] = 1073741824; /* Possibly out of bounds */  
    return s.d;  
}
```

fun(0)	⌘	3.14
fun(1)	⌘	3.14
fun(2)	⌘	3.1399998664856
fun(3)	⌘	2.00000061035156
fun(4)	⌘	3.14
fun(6)	⌘	Segmentation fault

- Result is system specific

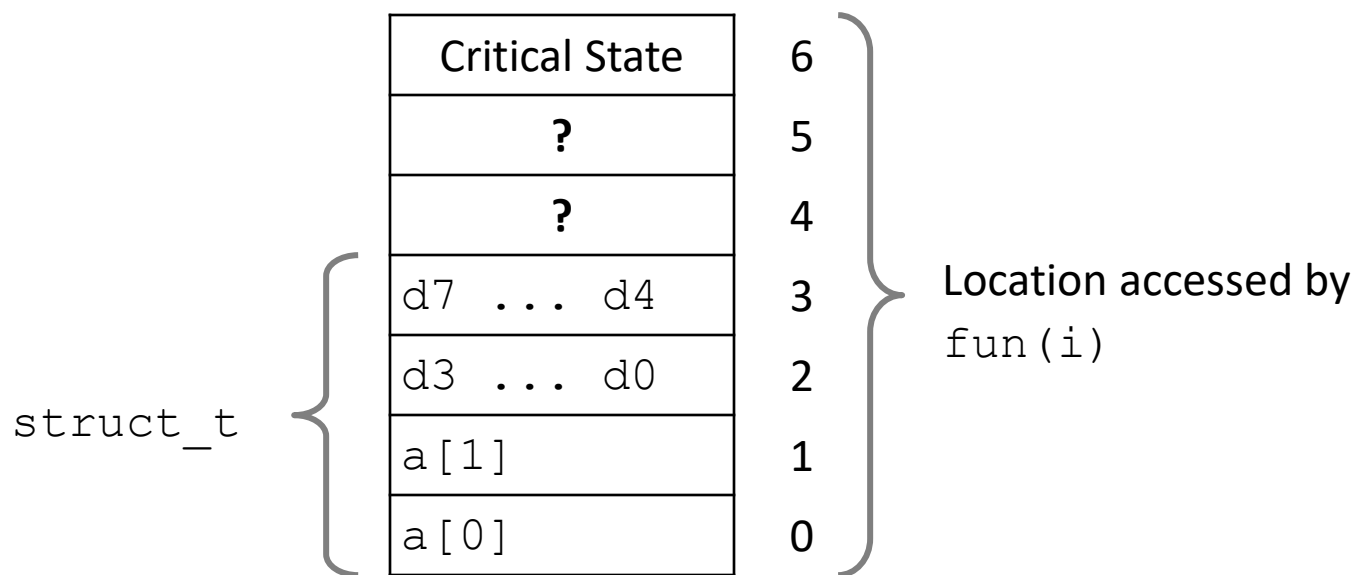


# Memory Referencing Bug Example

```
typedef struct {
    int a[2];
    double d;
} struct_t;
```

fun(0)	⌘	3.14
fun(1)	⌘	3.14
fun(2)	⌘	3.1399998664856
fun(3)	⌘	2.00000061035156
fun(4)	⌘	3.14
fun(6)	⌘	Segmentation fault

## Explanation:



# Memory Referencing Errors

## ■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

## ■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
  - Corrupted object logically unrelated to one being accessed
  - Effect of bug may be first observed long after it is generated

## ■ How can I deal with this?

- Program in Java, Ruby, Python, ML, ...
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

# Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
  - Easily see 10:1 performance range depending on how code written
  - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
  - How programs compiled and executed
  - How to measure program performance and identify bottlenecks
  - How to improve performance without destroying code modularity and generality

# Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```


**4.3ms****2.0 GHz Intel Core i7 Haswell**

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```

**81.8ms**

- Hierarchical memory organization
- Performance depends on access patterns
  - Including how step through multi-dimensional array

# Why The Performance Differs

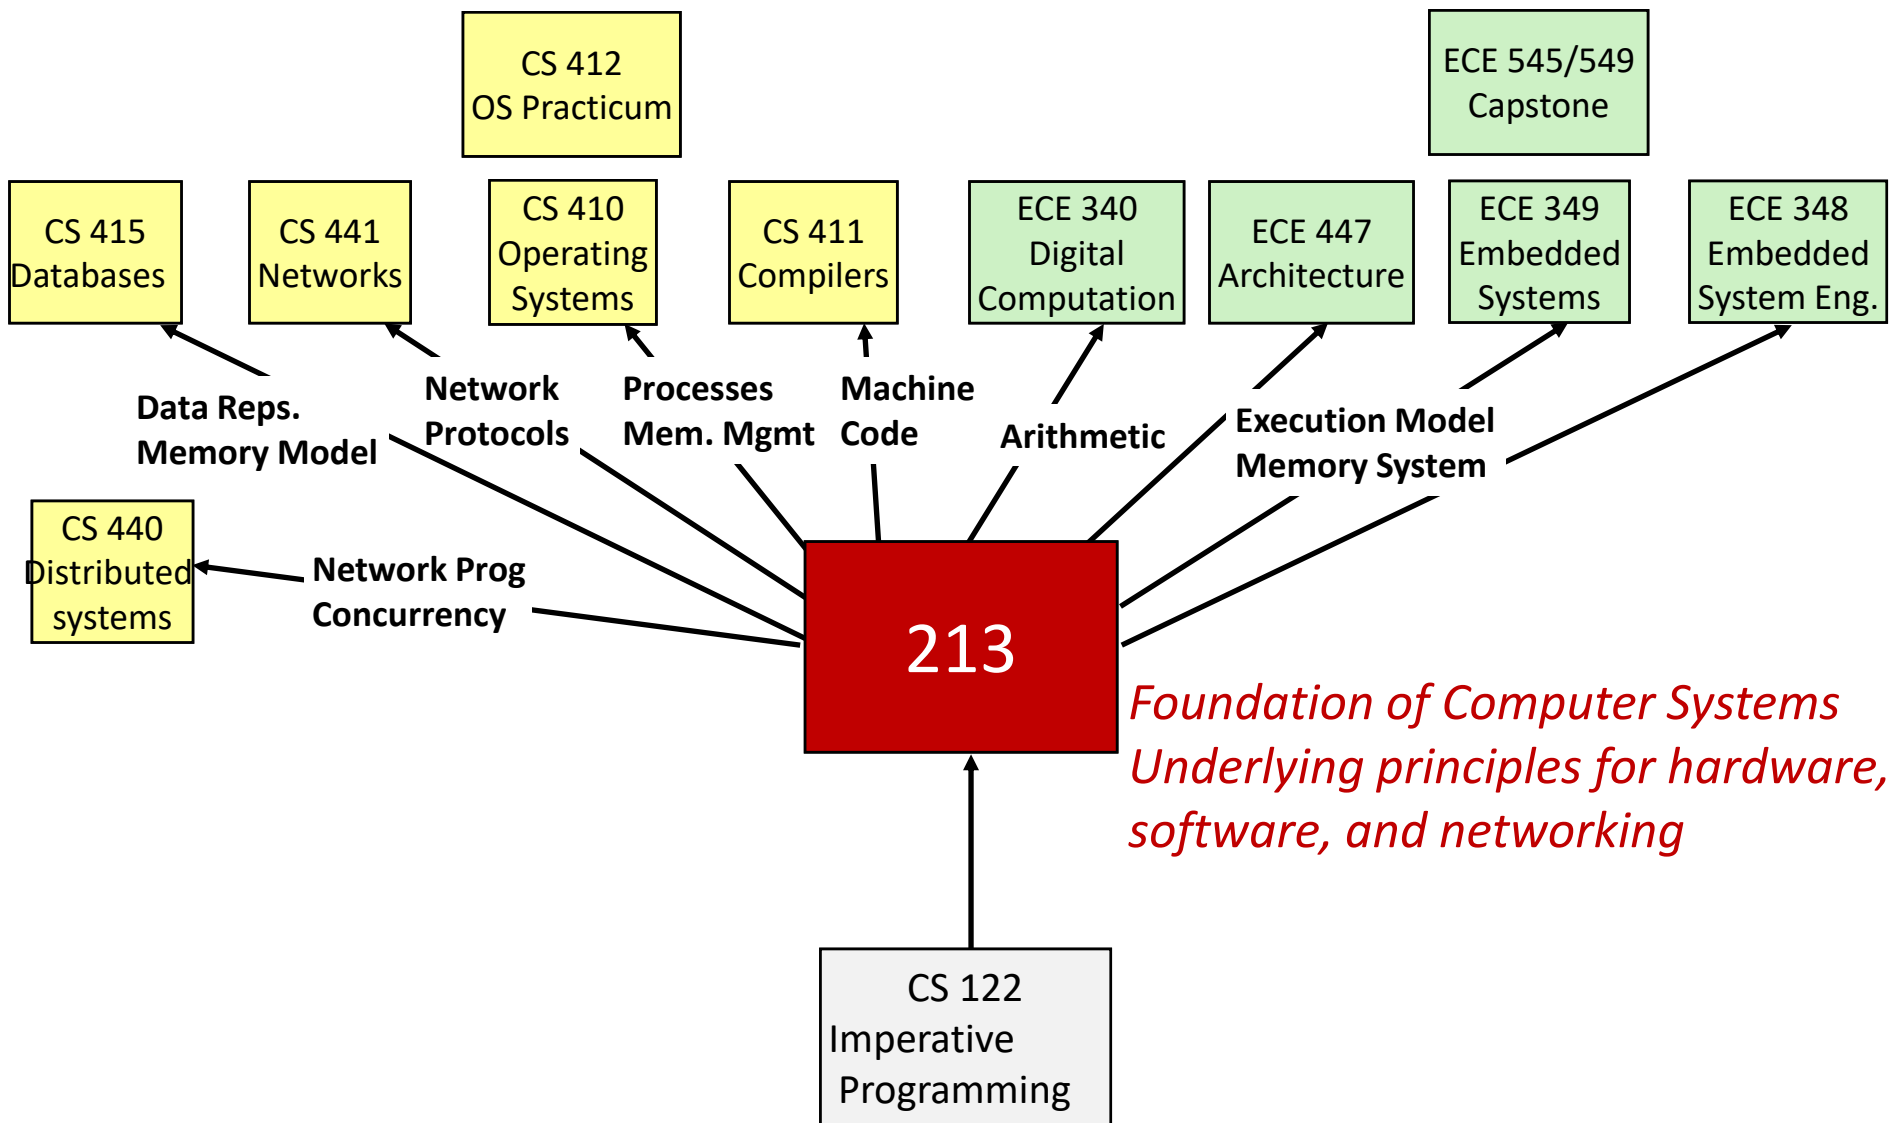


# Great Reality #5:

## Computers do more than execute programs

- **They need to get data in and out**
  - I/O system critical to program reliability and performance
  
- **They communicate with each other over networks**
  - Many system-level issues arise in presence of network
    - Concurrent operations by autonomous processes
    - Coping with unreliable media
    - Cross platform compatibility
    - Complex performance issues

# Role within CS/ECE Curriculum



# Course Perspective

## ■ Most Systems Courses are Builder-Centric

- Computer Architecture
  - Design pipelined processor in Verilog
- Operating Systems
  - Implement sample portions of operating system
- Compilers
  - Write compiler for simple language
- Networking
  - Implement and simulate network protocols



# Instructors

Randy Bryant



Dave O'Hallaron



# Cheating: Description

- **Please pay close attention, especially if this is your first semester at CMU**
- **What is cheating?**
  - Sharing code: by copying, retyping, **looking at**, or supplying a file
  - Describing: verbal description of code from one person to another.
  - Coaching: helping your friend to write a lab, line by line
  - Searching the Web for solutions
  - Copying code from a previous course or online solution
    - You are only allowed to use code we supply, or from the CS:APP website
- **What is NOT cheating?**
  - Explaining how to use systems or tools
  - Helping others with high-level design issues
- **See the course syllabus for details.**
  - Ignorance is not an excuse

# Cheating: Consequences

## ■ Penalty for cheating:

- Removal from course with failing grade (no exceptions!)
- Permanent mark on your record
- Your instructors' personal contempt

## ■ Detection of cheating:

- We have sophisticated tools for detecting code plagiarism
- Last Fall, 20 students were caught cheating and failed the course.
- Some were expelled from the University

## ■ Don't do it!

- Start early
- Ask the staff for help when you get stuck

# Textbooks

## ■ Randal E. Bryant and David R. O'Hallaron,

- *Computer Systems: A Programmer's Perspective*, **Third Edition** (CS:APP3e), Pearson, 2016
- <http://csapp.cs.cmu.edu>
- This book really matters for the course!
  - How to solve labs
  - Practice problems typical of exam problems

## ■ Brian Kernighan and Dennis Ritchie,

- *The C Programming Language*, Second Edition, Prentice Hall, 1988
- Still the best book about C, from the originators

# Course Components

## ■ Lectures

- Higher level concepts

## ■ Recitations

- Applied concepts, important tools and skills for labs, clarification of lectures, exam coverage

## ■ Labs (7)

- The heart of the course
- 1-2 weeks each
- Provide in-depth understanding of an aspect of systems
- Programming and measurement

## ■ Exams (midterm + final)

- Test your understanding of concepts & mathematical principles

# Getting Help

## ■ Class Web page: <http://www.cs.cmu.edu/~213>

- Complete schedule of lectures, exams, and assignments
- Copies of lectures, assignments, exams, solutions
- Clarifications to assignments

## ■ Blackboard and Piazza

- We won't be using Blackboard or Piazza for the course

# Getting Help

## ■ Staff mailing list: [15-213-staff@cs.cmu.edu](mailto:15-213-staff@cs.cmu.edu)

- Use this for all communication with the teaching staff
- Always CC staff mailing list during email exchanges
- Send email to individual instructors only to schedule appointments

## ■ Office hours (starting Tue Sept 3):

- SMTWR, 5:30-7:30pm, WeH 5207

## ■ 1:1 Appointments

- You can schedule 1:1 appointments with any of the teaching staff

# Policies: Labs And Exams

## ■ Work groups

- You must work alone on all lab assignments

## ■ Handins

- Labs due at 11:59pm on Tues or Thurs
- Electronic handins using **Autolab** (no exceptions!)

## ■ Exams

- Exams will be online in network-isolated clusters
- Held over multiple days. Self-scheduled; just sign up!

## ■ Appealing grades

- In **writing** to Prof O'Hallaron within 7 days of completion of grading
- Follow formal procedure described in syllabus



# Facilities

## ■ Labs will use the Intel Computer Systems Cluster

- The “shark machines”
- `linux> ssh shark.ics.cs.cmu.edu`
- 21 servers donated by Intel for 213
  - 10 student machines (for student logins)
  - 1 head node (for Autolab server and instructor logins)
  - 10 grading machines (for autograding)
- Each server: iCore 7: 8 Nehalem cores, 32 GB DRAM, RHEL 6.1
- Rack-mounted in Gates machine room
- Login using your Andrew ID and password
- **Getting help with the cluster machines:**
  - Please direct questions to staff mailing list

# Timeliness

## ■ Grace days

- **5 grace days** for the semester
- **Limit of 2 grace days per lab used automatically**
- Covers scheduling crunch, out-of-town trips, illnesses, minor setbacks
- Save them until late in the term!

## ■ Lateness penalties

- Once grace day(s) used up, get penalized **15% per day**
- No handins later than **3 days after due date**

## ■ Catastrophic events

- Major illness, death in family, ...
- Formulate a plan (with your academic advisor) to get back on track

## ■ Advice

- Once you start running late, it's really hard to catch up

# Other Rules of the Lecture Hall

- Laptops: permitted
- Electronic communications: *forbidden*
  - No email, instant messaging, cell phone calls, etc
- Presence in lectures, recitations: voluntary, recommended
- No recordings of ANY KIND

# Policies: Grading

- **Exams (50%): midterm (20%), final (30%)**
- **Labs (50%): weighted according to effort**
- **Final grades based on a straight scale.**

# Programs and Data

## ■ Topics

- Bits operations, arithmetic, assembly language programs
- Representation of C control and data structures
- Includes aspects of architecture and compilers

## ■ Assignments

- L1 (datalab): Manipulating bits
- L2 (bomblab): Defusing a binary bomb
- L3 (attacklab): The basics of code injection attacks

# The Memory Hierarchy

## ■ Topics

- Memory technology, memory hierarchy, caches, disks, locality
- Includes aspects of architecture and OS

## ■ Assignments

- L4 (cachelab): Building a cache simulator and optimizing for locality.
  - Learn how to exploit locality in your programs.

# Exceptional Control Flow

## ■ Topics

- Hardware exceptions, processes, process control, Unix signals, nonlocal jumps
- Includes aspects of compilers, OS, and architecture

## ■ Assignments

- L5 (tshlab): Writing your own Unix shell.
  - A first introduction to concurrency

# Virtual Memory

## ■ Topics

- Virtual memory, address translation, dynamic storage allocation
- Includes aspects of architecture and OS

## ■ Assignments

- L6 (malloclab): Writing your own malloc package
  - Get a real feel for systems-level programming



# Networking, and Concurrency

## ■ Topics

- High level and low-level I/O, network programming
- Internet services, Web servers
- concurrency, concurrent server design, threads
- I/O multiplexing with select
- Includes aspects of networking, OS, and architecture

## ■ Assignments

- L7 (proxylab): Writing your own Web proxy
  - Learn network programming and more about concurrency and synchronization.

# Lab Rationale

- **Each lab has a well-defined goal such as solving a puzzle or winning a contest**
- **Doing the lab should result in new skills and concepts**
- **We try to use competition in a fun and healthy way**
  - Set a reasonable threshold for full credit
  - Post intermediate results (anonymized) on Autolab scoreboard for glory!

# Autolab (<https://autolab.cs.cmu.edu>)

## ■ Labs are provided by the CMU Autolab system

- Project page: <http://autolab.cs.cmu.edu>
- Developed by CMU faculty and students
- Key ideas: Autograding and Scoreboards
  - **Autograding:** Providing you with instant feedback.
  - **Scoreboards:** Real-time, rank-ordered, and anonymous summary.
- Used by over 3,000 students each semester

## ■ With Autolab you can use your Web browser to:

- Download the lab materials
- Handin your code for autograding by the Autolab server
- View the class scoreboard
- View the complete history of your code handins, autograded results, instructor's evaluations, and gradebook.
- View the TA annotations of your code for Style points.

# Autolab accounts

- **Students enrolled 10am on Mon, Aug 26 have Autolab accounts**
- **You must be enrolled to get an account**
  - Autolab is not tied in to the Hub's rosters
  - If you add in, contact [15-213-staff@cs.cmu.edu](mailto:15-213-staff@cs.cmu.edu) for an account
- **For those who are waiting to add in, the first lab (datalab) will be available on the Schedule page of the course Web site.**

# Waitlist questions

- 15-213: Catherine Fichtner ([cathyf@cs.cmu.edu](mailto:cathyf@cs.cmu.edu))
- 18-213: Zara Collier ([zcollier@andrew.cmu.edu](mailto:zcollier@andrew.cmu.edu))
- 15-513: Catherine Fichtner ([cathyf@cs.cmu.edu](mailto:cathyf@cs.cmu.edu))
  
- Please don't contact the instructors with waitlist questions.

*Welcome  
and Enjoy!*